

PSTAT 131 Final Project

Alec Chen

2022-06-09

Contents

Introduction	1
Data Codebook	1
Data Preparation	2
Exploratory Data Analysis	2
Data Cleaning	3
Data Pre-processing	5
Lasso Regresion	13
SVM	15
Regression Tree	19
Random Forest	25

Introduction

I am interested in exploring what makes a song popular or not and what predictors contribute the most with EDA and thus based on that create 4 different machine learning models to predict songs' popularity. The data set consists of songs of different genres, artists, energy, liveliness, and other traits. Some of them top the chart list and become very popular and others are not so common among the crowd. I believe that each song has some unique combination of attributes that makes it popular. We are considering the data from Spotify, a popular application for listening to music. Each song has been rated on different factors in the data. With this analysis, a music company can predetermine how popular the song can come about to be. The model can also be used by companies like Spotify to predict the popularity of an upcoming song and thus suggest it to its user. To recommend new music to users, and to be able to internally classify songs, Spotify assigns each song value from 13 different features. These features are mostly numerical values but include some categorical data as well. Spotify also assigns each song a popularity score, based on the total number of clicks.

Data Collection: I found this data set from Kaggle <https://www.kaggle.com/datasets/zaheenhamidani/ultimate-spotify-tracks-db?resource=download> which consists of the features of songs. There are 10,000 songs per genre. There are 26 genres so it is a total of 232,725 tracks. This size of data set could be too big and inefficient to run models, therefore, I have randomly stratified the data with a size of 5000 to make sure the data set is still able to represent large data set.

Data Codebook

Numerical: -acousticness: A confidence measure from 0.0 to 1.0 of whether the track is acoustic. 1.0 represents high confidence the track is acoustic. -danceability: Danceability describes how suitable a track is for dancing based on a combination of musical elements including tempo, rhythm stability, beat strength, and overall regularity. A value of 0.0 is least danceable and 1.0 is most danceable. -duration_ms: The duration of the track in milliseconds. -energy: Energy is a measure from 0.0 to 1.0 and represents a perceptual measure of

intensity and activity. -liveliness: Detects the presence of an audience in the recording. -instrumentalness: Predicts whether a track contains no vocals. -loudness: The overall loudness of a track in decibels (dB) -speechiness: Speechiness detects the presence of spoken words in a track. -valence: A measure from 0.0 to 1.0 describing the musical positiveness conveyed by a track. -time_signature: An estimated time signature. -tempo: The overall estimated tempo of a track in beats per minute (BPM). Dummy Code: -mode: 0 = Minor, 1 = Major

Data Preparation

I loaded the necessary packages for later EDA and fit in linear regression, random forest, SVM for regression and boosted tree.

```
library(tidymodels)
library(tidyverse)
library(ISLR)
library(ISLR2)
library(discrim)
library(poissonreg)
library(corr)
library(ggplot2)
library(corrplot)
library(ggthemes)
library(kernlab)
library(e1071)
library(caret)
library(rpart)
library(rpart.plot)
tidymodels_prefer()

#Import the Data
spotify <- read_csv("Data/SpotifyFeatures.csv")
spotify
```

```
## # A tibble: 232,725 x 18
##   genre artist_name track_name track_id popularity acousticness danceability
##   <chr> <chr>         <chr>      <chr>      <dbl>      <dbl>      <dbl>
## 1 Movie Henri Salvador C'est bea~ OBRj06g~      0      0.611      0.389
## 2 Movie Martin & les ~ Perdu d'a~ OBjC1Nf~      1      0.246      0.59
## 3 Movie Joseph Willia~ Don't Let~ OCoSDzo~      3      0.952      0.663
## 4 Movie Henri Salvador Dis-moi M~ OGc6TVm~      0      0.703      0.24
## 5 Movie Fabien Nataf Ouverture OIuslXp~      4      0.95      0.331
## 6 Movie Henri Salvador Le petit ~ OMf1jKa~      0      0.749      0.578
## 7 Movie Martin & les ~ Premières~ ONUiKYR~      2      0.344      0.703
## 8 Movie Laura Mayne Let Me Le~ OPbIF9Y~     15      0.939      0.416
## 9 Movie Chorus Helka OST6uPf~      0      0.00104      0.734
## 10 Movie Le Club des J~ Les bisou~ OVSqZ3K~     10      0.319      0.598
## # ... with 232,715 more rows, and 11 more variables: duration_ms <dbl>,
## #   energy <dbl>, instrumentalness <dbl>, key <chr>, liveness <dbl>,
## #   loudness <dbl>, mode <chr>, speechiness <dbl>, tempo <dbl>,
## #   time_signature <chr>, valence <dbl>
```

Exploratory Data Analysis

Data Cleaning

We have 5,000 observations and 12 variables with two classes as numeric and integers dropping somer of the missing values rows and get a summary of the data set for better regression analysis.

```
# Data Set Summary
```

```
str(spotify)
```

```
## spec_tbl_df [232,725 x 18] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ genre          : chr [1:232725] "Movie" "Movie" "Movie" "Movie" ...
## $ artist_name    : chr [1:232725] "Henri Salvador" "Martin & les fées" "Joseph Williams" "Henri Sa
## $ track_name     : chr [1:232725] "C'est beau de faire un Show" "Perdu d'avance (par Gad Elmaleh)"
## $ track_id       : chr [1:232725] "0BRj06ga9RKCKjfdqeFgWV" "0BjC1NfoE00usryehmNudP" "0CoSDzoNIKCRs
## $ popularity     : num [1:232725] 0 1 3 0 4 0 2 15 0 10 ...
## $ acousticness   : num [1:232725] 0.611 0.246 0.952 0.703 0.95 0.749 0.344 0.939 0.00104 0.319 ...
## $ danceability   : num [1:232725] 0.389 0.59 0.663 0.24 0.331 0.578 0.703 0.416 0.734 0.598 ...
## $ duration_ms    : num [1:232725] 99373 137373 170267 152427 82625 ...
## $ energy         : num [1:232725] 0.91 0.737 0.131 0.326 0.225 0.0948 0.27 0.269 0.481 0.705 ...
## $ instrumentalness: num [1:232725] 0 0 0 0 0.123 0 0 0 0.00086 0.00125 ...
## $ key            : chr [1:232725] "C#" "F#" "C" "C#" ...
## $ liveness       : num [1:232725] 0.346 0.151 0.103 0.0985 0.202 0.107 0.105 0.113 0.0765 0.349 ..
## $ loudness       : num [1:232725] -1.83 -5.56 -13.88 -12.18 -21.15 ...
## $ mode           : chr [1:232725] "Major" "Minor" "Minor" "Major" ...
## $ speechiness    : num [1:232725] 0.0525 0.0868 0.0362 0.0395 0.0456 0.143 0.953 0.0286 0.046 0.02
## $ tempo          : num [1:232725] 167 174 99.5 171.8 140.6 ...
## $ time_signature : chr [1:232725] "4/4" "4/4" "5/4" "4/4" ...
## $ valence        : num [1:232725] 0.814 0.816 0.368 0.227 0.39 0.358 0.533 0.274 0.765 0.718 ...
## - attr(*, "spec")=
## .. cols(
## ..   genre = col_character(),
## ..   artist_name = col_character(),
## ..   track_name = col_character(),
## ..   track_id = col_character(),
## ..   popularity = col_double(),
## ..   acousticness = col_double(),
## ..   danceability = col_double(),
## ..   duration_ms = col_double(),
## ..   energy = col_double(),
## ..   instrumentalness = col_double(),
## ..   key = col_character(),
## ..   liveness = col_double(),
## ..   loudness = col_double(),
## ..   mode = col_character(),
## ..   speechiness = col_double(),
## ..   tempo = col_double(),
## ..   time_signature = col_character(),
## ..   valence = col_double()
## .. )
## - attr(*, "problems")=<externalptr>
```

```
colSums(is.na(spotify))
```

```
##           genre      artist_name      track_name      track_id
##           0           0           0           0
## popularity acousticness danceability duration_ms
##           0           0           0           0
```

```
##          energy instrumentalness          key          liveness
##              0              0              0              0
##          loudness              mode          speechiness          tempo
##              0              0              0              0
##    time_signature          valence
##              0              0
```

```
spotify <- drop_na(spotify)
summary(spotify)
```

```
##      genre          artist_name          track_name          track_id
## Length:232725    Length:232725    Length:232725    Length:232725
## Class :character  Class :character  Class :character  Class :character
## Mode  :character  Mode  :character  Mode  :character  Mode  :character
##
##
##      popularity          acousticness          danceability          duration_ms
## Min.   : 0.00    Min.   :0.00000    Min.   :0.0569    Min.   : 15387
## 1st Qu.: 29.00    1st Qu.:0.0376    1st Qu.:0.4350    1st Qu.: 182857
## Median : 43.00    Median :0.2320    Median :0.5710    Median : 220427
## Mean   : 41.13    Mean   :0.3686    Mean   :0.5544    Mean   : 235122
## 3rd Qu.: 55.00    3rd Qu.:0.7220    3rd Qu.:0.6920    3rd Qu.: 265768
## Max.   :100.00    Max.   :0.9960    Max.   :0.9890    Max.   :5552917
##      energy          instrumentalness          key          liveness
## Min.   :2.03e-05    Min.   :0.0000000    Length:232725    Min.   :0.00967
## 1st Qu.:3.85e-01    1st Qu.:0.0000000    Class :character  1st Qu.:0.09740
## Median :6.05e-01    Median :0.0000443    Mode  :character  Median :0.12800
## Mean   :5.71e-01    Mean   :0.1483012                    Mean   :0.21501
## 3rd Qu.:7.87e-01    3rd Qu.:0.0358000                    3rd Qu.:0.26400
## Max.   :9.99e-01    Max.   :0.9990000                    Max.   :1.00000
##      loudness          mode          speechiness          tempo
## Min.   : -52.457    Length:232725    Min.   :0.0222    Min.   : 30.38
## 1st Qu.: -11.771    Class :character  1st Qu.:0.0367    1st Qu.: 92.96
## Median : -7.762    Mode  :character  Median :0.0501    Median :115.78
## Mean   : -9.570                    Mean   :0.1208    Mean   :117.67
## 3rd Qu.: -5.501                    3rd Qu.:0.1050    3rd Qu.:139.05
## Max.   : 3.744                    Max.   :0.9670    Max.   :242.90
##    time_signature          valence
## Length:232725    Min.   :0.0000
## Class :character  1st Qu.:0.2370
## Mode  :character  Median :0.4440
##                      Mean   :0.4549
##                      3rd Qu.:0.6600
##                      Max.   :1.0000
```

With the unprocessed data set we need to rename some of the columns for better access, change the class of the explicit column from character to numeric and creates subset of data set for correlation matrices.

```
spotify<- spotify %>% select("danceability","energy","loudness","mode","speechiness",
                             "acousticness","instrumentalness","liveness","valence","tempo","duration_ms","popularity")
spotify <- spotify %>%
  mutate(mode = as.numeric(case_when(
    (mode == "Major") ~ "1",
    (mode == "Minor")~ "0")))
```

Data Pre-processing


We splitting the data set into two, a training set and a testing set and also perform a cross-fold validation.

```
# Data Train-Test Split
set.seed(111)
# Reduced data size for efficiency
spotify <- spotify[sample(1:nrow(spotify), 5000,
  replace=FALSE),]
# Split data set 80% training 20% testing
spotify_split <- initial_split(spotify, prop = 0.80,
  strata = popularity)
spotify_train <- training(spotify_split)
spotify_test <- testing(spotify_split)
# Cross fold validation with 10 folds
spotify_folds <- vfold_cv(spotify_train, v = 10)
spotify_folds


## # 10-fold cross-validation
## # A tibble: 10 x 2
##   splits          id
##   <list>         <chr>
## 1 <split [3600/400]> Fold01
## 2 <split [3600/400]> Fold02
## 3 <split [3600/400]> Fold03
## 4 <split [3600/400]> Fold04
## 5 <split [3600/400]> Fold05
## 6 <split [3600/400]> Fold06
## 7 <split [3600/400]> Fold07
## 8 <split [3600/400]> Fold08
## 9 <split [3600/400]> Fold09
## 10 <split [3600/400]> Fold10
```

With the correlation matrices of the features we find that most of the features have little or no correlation with one another beside the loudness and energy have a stronger positive correlation and acousticness and energy have a negative correlation. And, the histogram gives us the visualization of the popularity distribution within the data set. With Principal Components Regression we are able to see the intercept term in the test RMSE is 18.11 if we add in the first principal component it drops to 17.14. We can see that adding additional principal components actually leads to an increase in test RMSE. Thus, it appears that it would be optimal to only use two principal components in the final model. For the variance explained by using just the first principal component, we can explain 31.84% of the variation in the response variable. by adding in the second principal component, we can explain 46.21% of the variation in the response variable. We can explain more variance by using more principal components.

```
#Correlation of the predictors
corrplot::corrplot(cor(spotify), method = 'color')
```



```
# Data set song's popularity distribution
spotify %>%
  ggplot(aes(x = popularity)) +
  geom_histogram(bins = 60) +
  theme_bw()
```




```
#PCR
set.seed(111)
#fit PCR model
library(pls)
PCR_model <- pcr(popularity~., data=spotify, scale=TRUE, validation="CV")
summary(PCR_model)
```

```
## Data:      X dimension: 5000 11
## Y dimension: 5000 1
## Fit method: svdpc
## Number of components considered: 11
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           18.11   17.14   16.74   16.72   16.67   16.6   16.4
## adjCV        18.11   17.14   16.74   16.72   16.66   16.6   16.4
##      7 comps  8 comps  9 comps 10 comps 11 comps
## CV           16.21   16.04   15.87   15.88   15.85
## adjCV        16.21   16.04   15.87   15.87   15.85
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X           31.84   46.21   56.82   66.07   74.65   81.92   88.25
## popularity   10.57   14.71   14.99   15.60   16.45   18.39   20.36
##      8 comps  9 comps 10 comps 11 comps
## X           92.79   96.44   98.94   100.00
```


```
## popularity    22.02    23.67    23.67    23.92
```

```
#visualize cross-validation plots
```

```
validationplot(PCR_model)
```



```
validationplot(PCR_model, val.type="MSEP")
```



#

Model Building ## Linear Regression With the simple linear regression we can see that it predicts pretty poorly that the data points hardly follow a diagonal line. With the high of RMSE 15.68 and the low of rsq explaining 25% of the variance.

```
## Linear Regression
# Create a Recipe
spotify_recipe <- recipe(popularity ~ ., data = spotify_train) %>%
  step_dummy(all_nominal_predictors())

lm_model <- linear_reg() %>%
  set_engine("lm")
lm_wflow <- workflow() %>%
  add_model(lm_model) %>%
  add_recipe(spotify_recipe)

lm_fit <- fit(lm_wflow, spotify_train)
lm_fit %>%
  # This returns the parsnip object:
  extract_fit_parsnip() %>%
  # Now tidy the linear model object:
  tidy()
```

```
## # A tibble: 12 x 5
##   term                estimate std.error statistic  p.value
##   <chr>                <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)         58.0      2.62     22.2 8.09e-103
## 2 danceability        16.4      1.86      8.83 1.54e- 18
## 3 energy              -4.04      2.10     -1.92 5.44e- 2
## 4 loudness             0.675    0.0836     8.07 9.17e- 16
## 5 mode                -1.03      0.528    -1.95 5.08e- 2
## 6 speechiness         -7.44      1.74     -4.28 1.88e- 5
## 7 acousticness       -12.1      1.17    -10.4 5.61e- 25
## 8 instrumentality     -4.66      1.01     -4.61 4.11e- 6
## 9 liveness            -11.1      1.54     -7.16 9.63e- 13
## 10 valence            -15.8      1.26    -12.5 5.00e- 35
## 11 tempo              -0.0147    0.00857   -1.71 8.71e- 2
## 12 duration_ms         0.00000332 0.00000196    1.70 8.92e- 2
```

```
lm_fitt <- fit(lm_wflow, spotify_test)
lm_fitt %>%
  # This returns the parsnip object:
  extract_fit_parsnip() %>%
  # Now tidy the linear model object:
  tidy()
```

```
## # A tibble: 12 x 5
##   term                estimate std.error statistic  p.value
##   <chr>                <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)         58.3      5.23     11.1 2.81e-27
## 2 danceability        18.9      3.63      5.20 2.47e- 7
## 3 energy              -8.79      4.20     -2.09 3.65e- 2
## 4 loudness             0.759    0.162     4.69 3.19e- 6
## 5 mode                -2.27      1.06     -2.15 3.19e- 2
## 6 speechiness         -5.59      3.37     -1.66 9.74e- 2
## 7 acousticness       -12.4      2.36     -5.27 1.72e- 7
```

```
## 8 instrumentality -6.15      1.97      -3.12 1.87e- 3
## 9 liveness       -8.05      3.06      -2.63 8.65e- 3
## 10 valence       -12.1      2.50      -4.82 1.66e- 6
## 11 tempo         -0.0171    0.0174     -0.985 3.25e- 1
## 12 duration_ms   0.00000614 0.00000498    1.23 2.18e- 1
```

Train Recipe

```
spotify_train_res <- predict(lm_fit, new_data = spotify_train %>% select(-popularity))
spotify_train_res %>%
  head()
```


```
## # A tibble: 6 x 1
##   .pred
##   <dbl>
## 1  42.9
## 2  41.1
## 3  47.9
## 4  42.4
## 5  48.9
## 6  44.0
```

Test Recipe

```
spotify_train_res <- bind_cols(spotify_train_res, spotify_train %>% select(popularity))
spotify_train_res %>%
  head()
```

```
## # A tibble: 6 x 2
##   .pred popularity
##   <dbl>      <dbl>
## 1  42.9        26
## 2  41.1        19
## 3  47.9        27
## 4  42.4        29
## 5  48.9        24
## 6  44.0        28
```

```
spotify_train_res %>%
  ggplot(aes(x = .pred, y = popularity)) +
  geom_point(alpha = 0.2) +
  geom_abline(lty = 2) +
  theme_bw() +
  coord_obs_pred()
```



```
rmse(spotify_train_res, truth = popularity, estimate = .pred)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard      15.8
```

```
spotify_metrics <- metric_set(rmse, rsq, mae)
spotify_metrics(spotify_train_res, truth = popularity,
                estimate = .pred)
```

```
## # A tibble: 3 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard      15.8
## 2 rsq     standard       0.236
## 3 mae     standard      12.7
```

```
spotify_test_res <- predict(lm_fitt, new_data = spotify_test %>% select(-popularity))
spotify_test_res %>%
  head()
```


```
## # A tibble: 6 x 1
##   .pred
##   <dbl>
## 1  47.9
## 2  41.6
## 3  41.2
## 4  49.6
```

```
## 5 36.1
## 6 48.5

spotify_test_res <- bind_cols(spotify_test_res, spotify_test %>% select(popularity))
spotify_test_res %>%
  head()
```

```
## # A tibble: 6 x 2
##   .pred popularity
##   <dbl>      <dbl>
## 1  47.9         37
## 2  41.6         32
## 3  41.2         43
## 4  49.6         72
## 5  36.1         40
## 6  48.5         47
```

```
spotify_test_res %>%
  ggplot(aes(x = .pred, y = popularity)) +
  geom_point(alpha = 0.2) +
  geom_abline(lty = 2) +
  theme_bw() +
  coord_obs_pred()
```



```
rmse(spotify_test_res, truth = popularity, estimate = .pred)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 rmse    rmse         10.5
```

```
## 1 rmse      standard      15.7
spotify_metrics <- metric_set(rmse, rsq, mae)
spotify_metrics(spotify_test_res, truth = popularity,
                estimate = .pred)
```

```
## # A tibble: 3 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 rmse    standard      15.7
## 2 rsq     standard       0.259
## 3 mae     standard      12.5
```

Lasso Regression


To make sure that the linear regression might not be the best for predicting this data set, we use lasso regression to see if it provides a better prediction accuracy. In comparison with the simple linear regression we still see that it predicts pretty poorly with the high of RMSE 15.76 and the low of rsq explaining 25% of the variance.

```
# Lasso Regression
lasso_recipe <-
  recipe(formula = popularity ~ ., data = spotify_train) %>%
  step_dummy(all_nominal_predictors()) %>%
  step_novel(all_nominal_predictors()) %>%
  step_zv(all_predictors()) %>%
  step_normalize(all_predictors())

lasso_spec <-
  linear_reg(penalty = tune(), mixture = 1) %>%
  set_mode("regression") %>%
  set_engine("glmnet")

lasso_workflow <- workflow() %>%
  add_recipe(lasso_recipe) %>%
  add_model(lasso_spec)
penalty_grid <- grid_regular(penalty(range = c(-2, 2)), levels = 50)

tune_res <- tune_grid(
  lasso_workflow,
  resamples = spotify_folds,
  grid = penalty_grid
)
autoplot(tune_res)
```



```
collect_metrics(tune_res)
```

```
## # A tibble: 100 x 7
##   penalty .metric .estimator   mean     n std_err .config
##   <dbl> <chr>   <chr>     <dbl> <int>   <dbl> <chr>
## 1 0.01    rmse     standard  15.9     10 0.178 Preprocessor1_Model01
## 2 0.01    rsq      standard   0.233     10 0.00984 Preprocessor1_Model01
## 3 0.0121  rmse     standard  15.9     10 0.178 Preprocessor1_Model02
## 4 0.0121  rsq      standard   0.233     10 0.00984 Preprocessor1_Model02
## 5 0.0146  rmse     standard  15.9     10 0.178 Preprocessor1_Model03
## 6 0.0146  rsq      standard   0.233     10 0.00984 Preprocessor1_Model03
## 7 0.0176  rmse     standard  15.9     10 0.178 Preprocessor1_Model04
## 8 0.0176  rsq      standard   0.233     10 0.00984 Preprocessor1_Model04
## 9 0.0212  rmse     standard  15.9     10 0.178 Preprocessor1_Model05
## 10 0.0212  rsq      standard   0.233     10 0.00983 Preprocessor1_Model05
## # ... with 90 more rows
```

```
best_penalty <- select_best(tune_res, metric = "rsq")
best_penalty
```

```
## # A tibble: 1 x 2
##   penalty .config
##   <dbl> <chr>
## 1 0.0146 Preprocessor1_Model03
```

```
lasso_final <- finalize_workflow(lasso_workflow, best_penalty)
```

```
lasso_final_fit <- fit(lasso_final, data = spotify_train)
```

```
augment(lasso_final_fit, new_data = spotify_test) %>%
  rmse(truth = popularity, estimate = .pred)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard      15.8
```

```
augment(lasso_final_fit, new_data = spotify_test) %>%
  rsq(truth = popularity, estimate = .pred)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rsq     standard      0.251
```

```
augment(lasso_final_fit, new_data = spotify_test) %>%
  rmse(truth = popularity, estimate = .pred)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard      15.8
```

SVM

With SVM models we start off with just fitting into radial model the RMSE was high as 15 and with the hyper parameters tuning of scaling and centering the rmse slightly decreases then with the cross validation has a more significant decreases on RMSE.

```
#SVM
#Preprocessing Model1
modelsvm = svm(spotify$popularity~.,spotify)
set.seed(1)
modell1 <- train(
  popularity~.,
  data = spotify,
  method = 'svmRadial'
)
modell1
```

```
## Support Vector Machines with Radial Basis Function Kernel
##
## 5000 samples
## 11 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 5000, 5000, 5000, 5000, 5000, ...
## Resampling results across tuning parameters:
##
##   C      RMSE      Rsquared   MAE
## 0.25  15.05170  0.3168386  11.81862
## 0.50  15.07151  0.3167673  11.82390
## 1.00  15.14422  0.3132502  11.87142
##
```

```
## Tuning parameter 'sigma' was held constant at a value of 0.08594709
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were sigma = 0.08594709 and C = 0.25.
```

#Preprocessing Model2

```
set.seed(1)

model2 <- train(
  popularity~.,
  data = spotify,
  method = 'svmRadial',
  preProcess = c("center", "scale")
)
model2
```

```
## Support Vector Machines with Radial Basis Function Kernel
```

```
##
```

```
## 5000 samples
```

```
## 11 predictor
```

```
##
```

```
## Pre-processing: centered (11), scaled (11)
```

```
## Resampling: Bootstrapped (25 reps)
```

```
## Summary of sample sizes: 5000, 5000, 5000, 5000, 5000, 5000, ...
```

```
## Resampling results across tuning parameters:
```

```
##
```

##	C	RMSE	Rsquared	MAE
##	0.25	15.05170	0.3168386	11.81862
##	0.50	15.07151	0.3167673	11.82390
##	1.00	15.14422	0.3132502	11.87142

```
##
```

```
## Tuning parameter 'sigma' was held constant at a value of 0.08594709
```

```
## RMSE was used to select the optimal model using the smallest value.
```

```
## The final values used for the model were sigma = 0.08594709 and C = 0.25.
```

Splitting data

```
set.seed(1)

inTraining <- createDataPartition(spotify$popularity, p = .80, list = FALSE)
training <- spotify[inTraining,]
testing <- spotify[-inTraining,]
set.seed(1)

model3 <- train(
  popularity ~ .,
  data = training,
  method = 'svmRadial',
  preProcess = c("center", "scale")
)
model3
```

```
## Support Vector Machines with Radial Basis Function Kernel
```

```
##
```

```
## 4000 samples
```

```
## 11 predictor
```

```
##
```

```
## Pre-processing: centered (11), scaled (11)
```



```

## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 4000, 4000, 4000, 4000, 4000, 4000, ...
## Resampling results across tuning parameters:
##
##      C      RMSE      Rsquared    MAE
##  0.25  14.97092  0.3159534  11.74475
##  0.50  15.00488  0.3150187  11.77565
##  1.00  15.10408  0.3096636  11.86376
##
## Tuning parameter 'sigma' was held constant at a value of 0.08991346
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were sigma = 0.08991346 and C = 0.25.

# calculate the RMSE and r2 to compare to the model above.
test.features = subset(testing, select=-c(popularity))
test.target = subset(testing, select=popularity)[,1]

predictions = predict(model3, newdata = test.features)

# RMSE
sqrt(mean((test.target - predictions)^2))

## [1] NA

# R2
cor(test.target, predictions) ^ 2

##           [,1]
## popularity 0.3162256

#Cross Validation
set.seed(1)
ctrl <- trainControl(
  method = "cv",
  number = 10,
)

set.seed(1)

model4 <- train(
  popularity ~ .,
  data = training,
  method = 'svmRadial',
  preProcess = c("center", "scale"),
  trCtrl = ctrl
)
model4

## Support Vector Machines with Radial Basis Function Kernel
##
## 4000 samples
## 11 predictor
##
## Pre-processing: centered (11), scaled (11)
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 4000, 4000, 4000, 4000, 4000, 4000, ...
## Resampling results across tuning parameters:

```

```

##
##      C      RMSE      Rsquared    MAE
##    0.25  14.97092  0.3159534  11.74475
##    0.50  15.00488  0.3150187  11.77565
##    1.00  15.10408  0.3096636  11.86376
##
## Tuning parameter 'sigma' was held constant at a value of 0.08991346
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were sigma = 0.08991346 and C = 0.25.
# calculate the RMSE and r2 to compare to the model above.
test.features = subset(testing, select=-c(popularity))
test.target = subset(testing, select=popularity)[,1]

predictions = predict(model4, newdata = test.features)

# RMSE
sqrt(mean((test.target - predictions)^2))

## [1] NA

# R2
cor(test.target, predictions) ^ 2

##                [,1]
## popularity 0.3162256
#Tuning Hyper Parameters
set.seed(1)

tuneGrid <- expand.grid(
  C = c(0.25, .5, 1),
  sigma = 0.1
)


model5 <- train(
  popularity ~ .,
  data = training,
  method = 'svmRadial',
  preProcess = c("center", "scale"),
  trControl = ctrl,
  tuneGrid = tuneGrid
)
model5

## Support Vector Machines with Radial Basis Function Kernel
##
## 4000 samples
## 11 predictor
##
## Pre-processing: centered (11), scaled (11)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 3600, 3601, 3601, 3600, 3599, 3601, ...
## Resampling results across tuning parameters:
##
##      C      RMSE      Rsquared    MAE
##    0.25  14.83539  0.3312740  11.63474

```

```
## 0.50 14.79984 0.3347396 11.60035
## 1.00 14.83162 0.3331954 11.63081
##
## Tuning parameter 'sigma' was held constant at a value of 0.1
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were sigma = 0.1 and C = 0.5.
```

```
plot(model5)
```

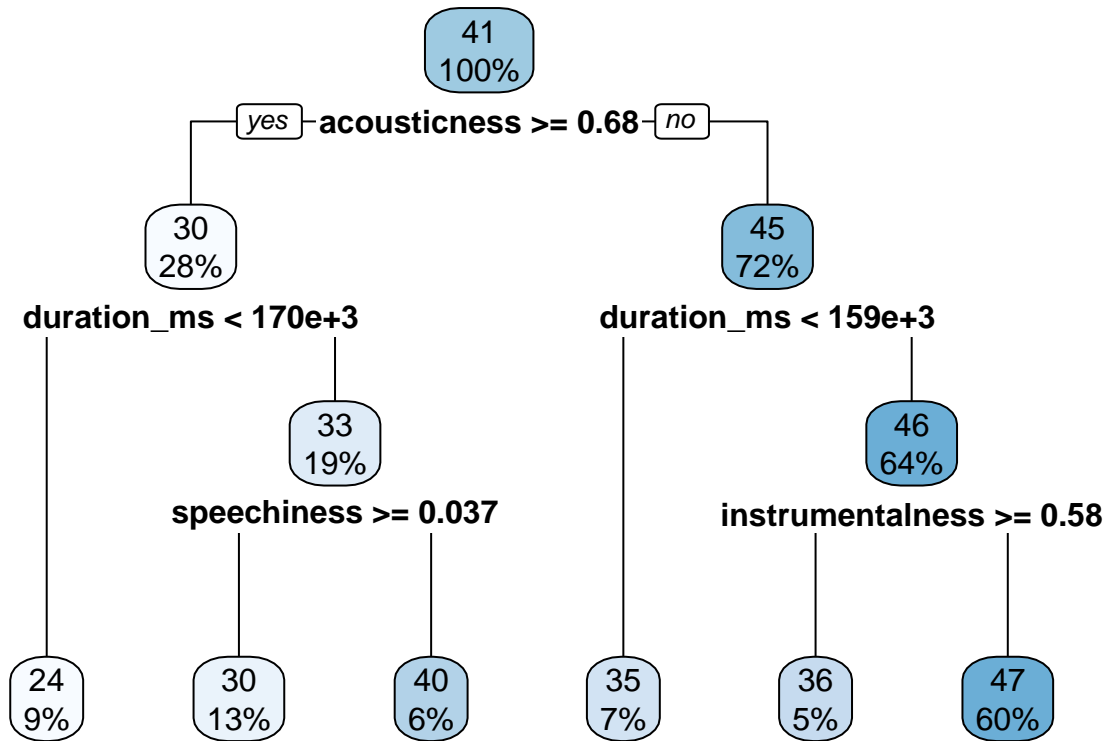


Regression Tree


```
#M1
tree_ml1 <- rpart(popularity ~ .,
                  method = "anova", data = spotify_train)
tree_ml1
```

```
## n= 4000
##
## node), split, n, deviance, yval
##      * denotes terminal node
##
## 1) root 4000 1308097.00 40.94250
##    2) acousticness>=0.6795 1133 346563.30 30.10591
##      4) duration_ms< 170300 375 95025.80 23.73067 *
##      5) duration_ms>=170300 758 228755.80 33.25989
##        10) speechiness>=0.03665 529 154602.00 30.30435 *
##        11) speechiness< 0.03665 229 58858.25 40.08734 *
##    3) acousticness< 0.6795 2867 775903.90 45.22497
##      6) duration_ms< 158539 288 117154.70 34.69792 *
##      7) duration_ms>=158539 2579 623269.20 46.40054
```

```
##      14) instrumentalness>=0.5845 189   25655.25 36.35450 *
##      15) instrumentalness< 0.5845 2390  577031.10 47.19498 *
rpart.plot(tree_ml1)
```




```
plotcp(tree_ml1)
```



```
#M2
tree_m12 <- rpart(
  formula = popularity ~ .,
  data    = spotify_train,
  method  = "anova",
  control = list(cp = 0, xval = 10)
)

plotcp(tree_m12)
abline(v = 12, lty = "dashed")
```

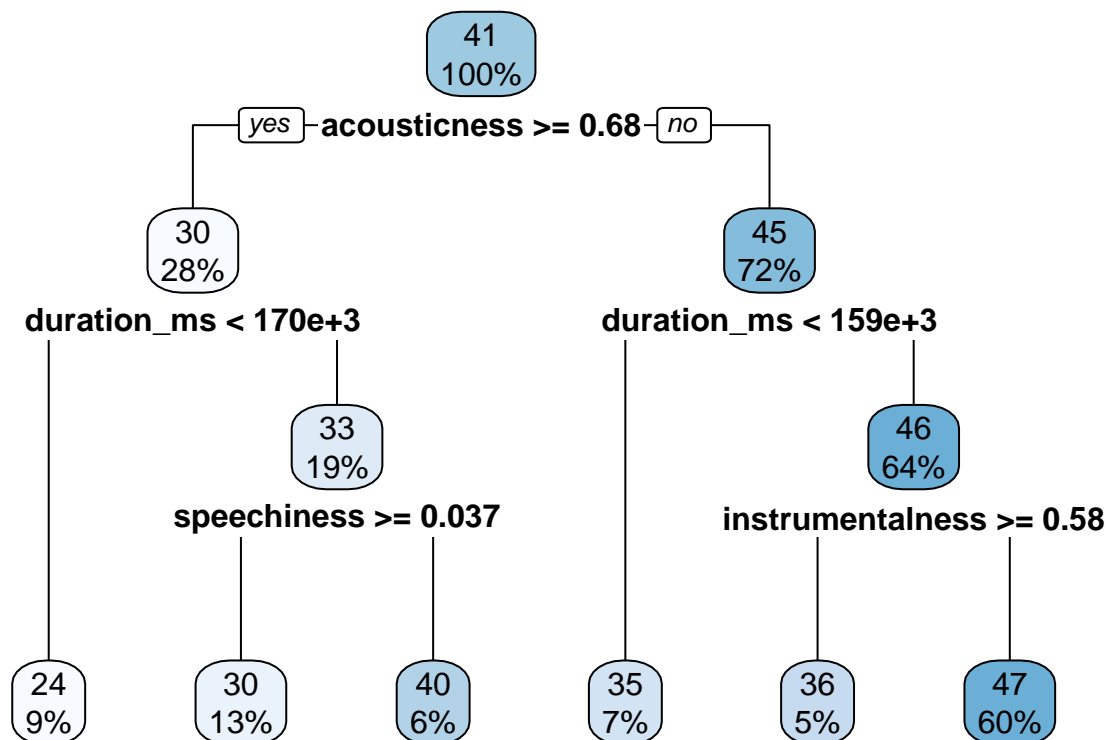


```
#M3
m3 <- rpart(
  formula = popularity ~ .,
  data     = spotify_train,
  method   = "anova",
  control  = list(minsplit = 10, maxdepth = 12, xval = 10)
)

tree_spec <- decision_tree() %>%
  set_engine("rpart")
reg_tree_spec <- tree_spec %>%
  set_mode("regression")
reg_tree_fit <- fit(reg_tree_spec, popularity ~ ., spotify_train)
augment(reg_tree_fit, new_data = spotify_test) %>%
  rmse(truth = popularity, estimate = .pred)

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 rmse    standard      16.2

reg_tree_fit %>%
  extract_fit_engine() %>%
  rpart.plot()
```




```

reg_tree_wf <- workflow() %>%
  add_model(reg_tree_spec %>% set_args(cost_complexity = tune())) %>%
  add_formula(popularity ~ .)

param_grid <- grid_regular(cost_complexity(range = c(-4, -1)), levels = 10)

tune_res <- tune_grid(
  reg_tree_wf,
  resamples = spotify_folds,
  grid = param_grid
)
autoplot(tune_res)

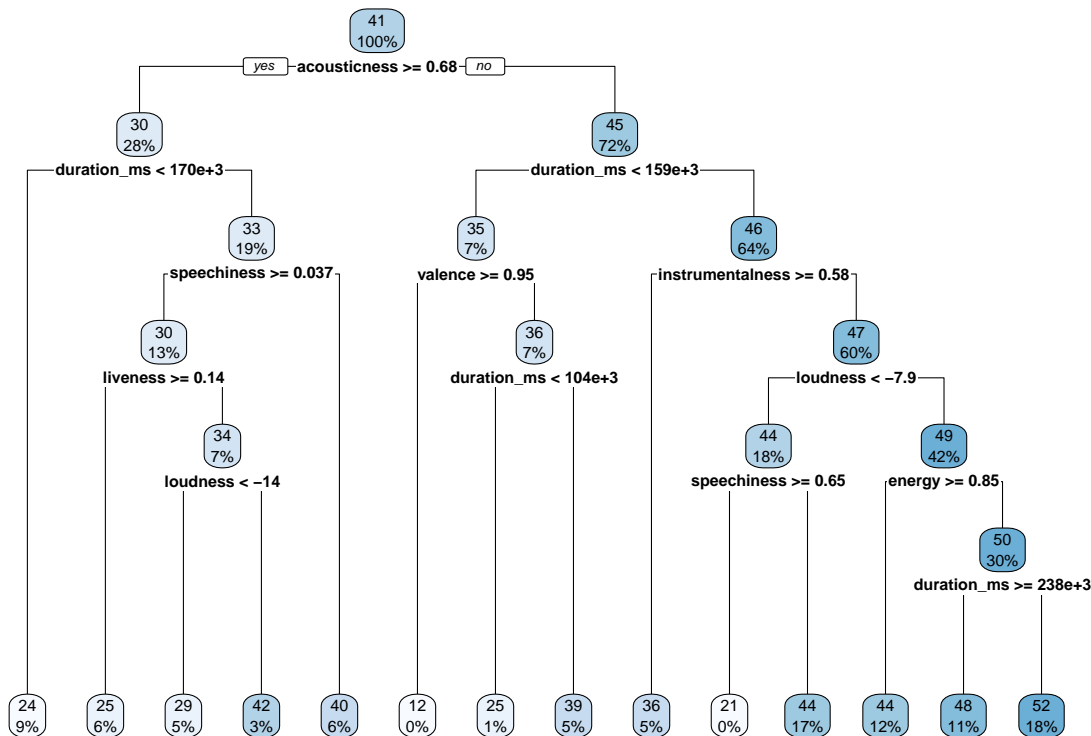
```



```
best_complexity <- select_best(tune_res, metric = "rmse")

reg_tree_final <- finalize_workflow(reg_tree_wf, best_complexity)

reg_tree_final_fit <- fit(reg_tree_final, data = spotify_train)
reg_tree_final_fit %>%
  extract_fit_engine() %>%
  rpart.plot()
```

```
augment(reg_tree_final_fit, new_data = spotify_test) %>%
  rmse(truth = popularity, estimate = .pred)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard      15.8
```

Random Forest


With boosted trees we can see that there are some predictors that weighted more than other predictors such as acousticness with high of importance of 15% incomparably with other predictors ranging 5-10%, which is inconsistent with what we have seen from the correlation matrix indicating that loudness is more correlated with popularity than other predictors. With the increase of tree depth doesnt help with the RMSE but it show the acousticness as taking more importance in the prediction.

```
bagging_spec <- rand_forest(mtry = .cols()) %>%
  set_engine("randomForest", importance = TRUE) %>%
  set_mode("regression")
bagging_fit <- fit(bagging_spec, popularity ~ .,
  data = spotify_train)
augment(bagging_fit, new_data = spotify_test) %>%
  rmse(truth = popularity, estimate = .pred)
```


```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard      14.6
```

```
augment(bagging_fit, new_data = spotify_test) %>%
  ggplot(aes(popularity, .pred)) +
  geom_abline() +
```

```
geom_point(alpha = 0.5)
```




```
library(vip)  
vip(bagging_fit)
```




```
rf_spec <- rand_forest(mtry = 6) %>%
  set_engine("randomForest", importance = TRUE) %>%
  set_mode("regression")
rf_fit <- fit(rf_spec, popularity ~ ., data = spotify_train)
augment(rf_fit, new_data = spotify_train) %>%
  rmse(truth = popularity, estimate = .pred)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard         6.07
```

```
augment(rf_fit, new_data = spotify_test) %>%
  ggplot(aes(popularity, .pred)) +
  geom_abline() +
  geom_point(alpha = 0.5)
```



```
vip(rf_fit)
```




Boosted Trees With boosted trees we can see that there are some predictors that weighted more than other predictors such as acousticness with high of importance of 15% incomparably with other predictors ranging 5-10%, which is inconsistent with what we have seen from the correlation matrix indicating that loudness is more correlated with popularity than other predictors. With the increase of tree depth doesnt help with the RMSE but it show the acousticness as taking more importnace in the prediction.


```
# Boosted Trees
boost_spec <- boost_tree(trees = 5000, tree_depth = 5) %>%
  set_engine("xgboost") %>%
  set_mode("regression")
boost_fit <- fit(boost_spec, popularity ~ ., data = spotify_train)
augment(boost_fit, new_data = spotify_test) %>%
  rmse(truth = popularity, estimate = .pred)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 rmse    standard      16.2
```

```
augment(boost_fit, new_data = spotify_test) %>%
  ggplot(aes(popularity, .pred)) +
  geom_abline() +
  geom_point(alpha = 0.5)
```



```
vip(boost_fit)
```



Conclusion Overall we have done linear regression, Ridge regression, Lasso regression, Random forest, SVM, and boosted tree all of them have the similar RMSE of around 16 on the testing set prediction beside the random forest has comparably the low of 6. Overall I am surprised by the high RMSE of the models that produce usually it would be best to be around 0.2 to 0.5. From the regression models fitting, we already see that the correlation among the predictors and the popularity are not linear thus they would have a poor prediction. High RMSE could be indicating the overfitting that it trains well with training set however does it poor prediction with the testing set. Usually Boosted trees should be performing better than random forest, however, the reason that the random forest performs better is because the overfitting of the boosted trees.